

session24_Assign2.R

Seshan

Tue Aug 28 07:47:35 2018

5. Problem Statement

1. Perform the below given activities:
 - a. Take a sample data set of your choice
 - b. Apply random forest, logistic regression using Spark R
 - c. Predict for new dataset

```
library(sparklyr)
library(ggplot2)
library(dplyr)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

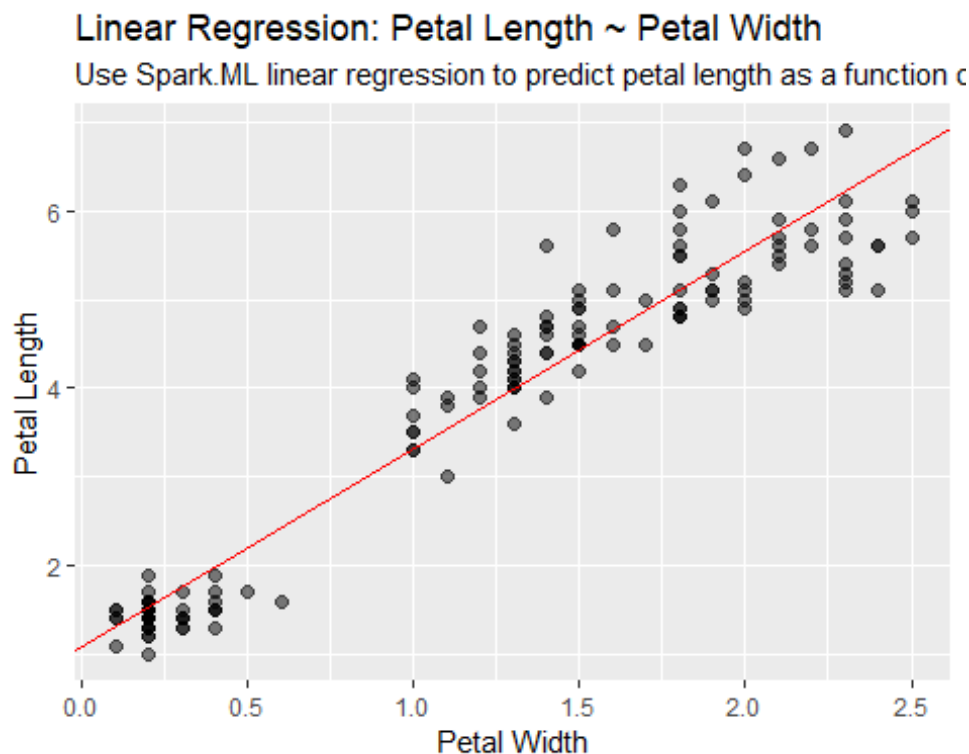
sc <- spark_connect(master = "local")
iris_tbl <- copy_to(sc, iris, "iris", overwrite = TRUE)
iris_tbl

## # Source:   table<iris> [?? x 5]
## # Database: spark_shell_connection
##   Sepal_Length Sepal_Width Petal_Length Petal_Width Species
##           <dbl>       <dbl>       <dbl>       <dbl> <chr>
## 1           5.1         3.5         1.4         0.2 setosa
## 2           4.9         3          1.4         0.2 setosa
## 3           4.7         3.2         1.3         0.2 setosa
## 4           4.6         3.1         1.5         0.2 setosa
## 5           5          3.6         1.4         0.2 setosa
## 6           5.4         3.9         1.7         0.4 setosa
## 7           4.6         3.4         1.4         0.3 setosa
## 8           5          3.4         1.5         0.2 setosa
## 9           4.4         2.9         1.4         0.2 setosa
```

```
## 10          4.9          3.1          1.5          0.1 setosa
## # ... with more rows

lm_model <- iris_tbl %>%
  select(Petal_Width, Petal_Length) %>%
  ml_linear_regression(Petal_Length ~ Petal_Width)

iris_tbl %>%
  select(Petal_Width, Petal_Length) %>%
  collect %>%
  ggplot(aes(Petal_Length, Petal_Width)) +
  geom_point(aes(Petal_Width, Petal_Length), size = 2, alpha = 0.5) +
  geom_abline(aes(slope = coef(lm_model)[["Petal_Width"]],
                  intercept = coef(lm_model)[["(Intercept)"]]),
              color = "red") +
  labs(
    x = "Petal Width",
    y = "Petal Length",
    title = "Linear Regression: Petal Length ~ Petal Width",
    subtitle = "Use Spark.ML linear regression to predict petal length as a function of petal width."
  )
```



```
pca_model <- tbl(sc, "iris") %>%
  select(-Species) %>%
  ml_pca()
print(pca_model)
```

```
## Explained variance:
##
##          PC1          PC2          PC3          PC4
## 0.924618723 0.053066483 0.017102610 0.005212184
##
## Rotation:
##          PC1          PC2          PC3          PC4
## Sepal_Length -0.36138659 -0.65658877 0.58202985 0.3154872
## Sepal_Width  0.08452251 -0.73016143 -0.59791083 -0.3197231
## Petal_Length -0.85667061 0.17337266 -0.07623608 -0.4798390
## Petal_Width  -0.35828920 0.07548102 -0.54583143 0.7536574
```

b. Apply random forest, logistic regression using Spark R

c. Predict for new dataset

```
#Random Forest
#Use Spark's Random Forest to perform regression or multiclass classification.

rf_model <- iris_tbl %>%
  ml_random_forest(Species ~ Petal_Length + Petal_Width, type = "classification")

rf_predict <- sdf_predict(rf_model, iris_tbl) %>%
  ft_string_indexer("Species", "Species_idx") %>%
  collect

## Warning in sdf_predict.ml_model(rf_model, iris_tbl): The signature
## sdf_predict(model, dataset) is deprecated and will be removed in a future
## version. Use sdf_predict(dataset, model) or ml_predict(model, dataset)
## instead.

table

## function (... , exclude = if (useNA == "no") c(NA, NaN), useNA = c("no",
##   "ifany", "always"), dnn = list.names(...), deparse.level = 1)
## {
##   list.names <- function(...) {
##     l <- as.list(substitute(list(...)))[-1L]
##     nm <- names(l)
##     fixup <- if (is.null(nm))
##       seq_along(l)
##     else nm == ""
##     dep <- vapply(l[fixup], function(x) switch(deparse.level +
##       1, "", if (is.symbol(x)) as.character(x) else "",
##       deparse(x, nlines = 1)[1L]), "")
##     if (is.null(nm))
```

```

##         dep
##     else {
##         nm[fixup] <- dep
##         nm
##     }
## }
## miss.use <- missing(useNA)
## miss.exc <- missing(exclude)
## useNA <- if (miss.use && !miss.exc && !match(NA, exclude,
##     nomatch = 0L))
##     "ifany"
## else match.arg(useNA)
## doNA <- useNA != "no"
## if (!miss.use && !miss.exc && doNA && match(NA, exclude,
##     nomatch = 0L))
##     warning("'exclude' containing NA and 'useNA' != \"no\" are a bit
contradicting")
## args <- list(...)
## if (!length(args))
##     stop("nothing to tabulate")
## if (length(args) == 1L && is.list(args[[1L]])) {
##     args <- args[[1L]]
##     if (length(dnn) != length(args))
##         dnn <- if (!is.null(argn <- names(args)))
##             argn
##         else paste(dnn[1L], seq_along(args), sep = ".")
## }
## bin <- 0L
## lens <- NULL
## dims <- integer()
## pd <- 1L
## dn <- NULL
## for (a in args) {
##     if (is.null(lens))
##         lens <- length(a)
##     else if (length(a) != lens)
##         stop("all arguments must have the same length")
##     fact.a <- is.factor(a)
##     if (doNA)
##         aNA <- anyNA(a)
##     if (!fact.a) {
##         a0 <- a
##         a <- factor(a, exclude = exclude)
##     }
##     add.na <- doNA
##     if (add.na) {
##         ifany <- (useNA == "ifany")
##         anNAc <- anyNA(a)
##         add.na <- if (!ifany || anNAc) {
##             ll <- levels(a)

```

```

##             if (add.ll <- !anyNA(ll)) {
##                 ll <- c(ll, NA)
##                 TRUE
##             }
##             else if (!ifany && !anNAc)
##                 FALSE
##             else TRUE
##         }
##     else FALSE
## }
## if (add.na)
##     a <- factor(a, levels = ll, exclude = NULL)
## else ll <- levels(a)
## a <- as.integer(a)
## if (fact.a && !miss.exc) {
##     ll <- ll[keep <- which(match(ll, exclude, nomatch = 0L) ==
##         0L)]
##     a <- match(a, keep)
## }
## else if (!fact.a && add.na) {
##     if (ifany && !aNA && add.ll) {
##         ll <- ll[!is.na(ll)]
##         is.na(a) <- match(a0, c(exclude, NA), nomatch = 0L) >
##             0L
##     }
##     else {
##         is.na(a) <- match(a0, exclude, nomatch = 0L) >
##             0L
##     }
## }
## nl <- length(ll)
## dims <- c(dims, nl)
## if (prod(dims) > .Machine$integer.max)
##     stop("attempt to make a table with >= 2^31 elements")
## dn <- c(dn, list(ll))
## bin <- bin + pd * (a - 1L)
## pd <- pd * nl
## }
## names(dn) <- dnn
## bin <- bin[!is.na(bin)]
## if (length(bin))
##     bin <- bin + 1L
## y <- array(tabulate(bin, pd), dims, dimnames = dn)
## class(y) <- "table"
## y
## }
## <bytecode: 0x0000000019782370>
## <environment: namespace:base>

```

```

partitions <- tbl(sc, "iris") %>%
  sdf_partition(training = 0.75, test = 0.25, seed = 1099)

fit <- partitions$training %>%
  ml_linear_regression(Petal_Length ~ Petal_Width)

estimate_mse <- function(df){
  sdf_predict(fit, df) %>%
    mutate(resid = Petal_Length - prediction) %>%
    summarize(mse = mean(resid ^ 2)) %>%
    collect
}

sapply(partitions, estimate_mse)

## Warning in sdf_predict.ml_model(fit, df): The signature sdf_predict(model,
## dataset) is deprecated and will be removed in a future version. Use
## sdf_predict(dataset, model) or ml_predict(model, dataset) instead.

## Warning: Missing values are always removed in SQL.
## Use `AVG(x, na.rm = TRUE)` to silence this warning

## Warning in sdf_predict.ml_model(fit, df): The signature sdf_predict(model,
## dataset) is deprecated and will be removed in a future version. Use
## sdf_predict(dataset, model) or ml_predict(model, dataset) instead.

## Warning: Missing values are always removed in SQL.
## Use `AVG(x, na.rm = TRUE)` to silence this warning

## $training.mse
## [1] 0.2374596
##
## $test.mse
## [1] 0.1898848

#Use ft_string_indexer and ft_index_to_string to convert a character column i
nto a numeric column and back again.

ft_string2idx <- iris_tbl %>%
  ft_string_indexer("Species", "Species_idx") %>%
  ft_index_to_string("Species_idx", "Species_remap") %>%
  collect

table(ft_string2idx$Species, ft_string2idx$Species_remap)

##
##          setosa versicolor virginica
## setosa          50             0         0
## versicolor       0             50         0
## virginica        0             0         50

```

```

ft_string2idx <- iris_tbl %>%
  sdf_mutate(Species_idx = ft_string_indexer(Species)) %>%
  sdf_mutate(Species_remap = ft_index_to_string(Species_idx)) %>%
  collect

ft_string2idx %>%
  select(Species, Species_idx, Species_remap) %>%
  distinct

## # A tibble: 3 x 3
##   Species      Species_idx Species_remap
##   <chr>          <dbl> <chr>
## 1 setosa              2 setosa
## 2 versicolor         0 versicolor
## 3 virginica           1 virginica

#Use Spark's Logistic regression to perform Logistic regression, modeling a binary outcome as a function of one or more explanatory variables.

# Prepare beaver dataset
beaver <- beaver2
beaver$activ <- factor(beaver$activ, labels = c("Non-Active", "Active"))
copy_to(sc, beaver, "beaver")

## # Source:   table<beaver> [?? x 4]
## # Database: spark_shell_connection
##    day  time  temp activ
##    <dbl> <dbl> <dbl> <chr>
##  1   307   930  36.6 Non-Active
##  2   307   940  36.7 Non-Active
##  3   307   950  36.9 Non-Active
##  4   307  1000  37.2 Non-Active
##  5   307  1010  37.2 Non-Active
##  6   307  1020  37.2 Non-Active
##  7   307  1030  37.2 Non-Active
##  8   307  1040  36.9 Non-Active
##  9   307  1050  37.0 Non-Active
## 10   307  1100  36.9 Non-Active
## # ... with more rows

beaver_tbl <- tbl(sc, "beaver")

glm_model <- beaver_tbl %>%
  mutate(binary_response = as.numeric(activ == "Active")) %>%
  ml_logistic_regression(binary_response ~ temp)

glm_model

## Formula: binary_response ~ temp
##
## Coefficients:

```

```
## (Intercept)      temp
##    550.52331    -14.69184

#First, we will copy the mtcars dataset into Spark.

mtcars_tbl <- copy_to(sc, mtcars, "mtcars")

# transform our data set, and then partition into 'training', 'test'
partitions <- mtcars_tbl %>%
  filter(hp >= 100) %>%
  sdf_mutate(cyl8 = ft_bucketizer(cyl, c(0,8,12))) %>%
  sdf_partition(training = 0.5, test = 0.5, seed = 888)

# fit a linear model to the training dataset
fit <- partitions$training %>%
  ml_linear_regression(mpg ~ wt + cyl)

# summarize the model
summary(fit)

## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.0947 -1.2747 -0.1129  1.0876  2.2185
##
## Coefficients:
## (Intercept)          wt          cyl
##   33.795576   -1.596247   -1.580360
##
## R-Squared: 0.8267
## Root Mean Squared Error: 1.437
```